

# Untitled

December 4, 2023

```
[68]: import pandas as pd
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
# import sys
# !{sys.executable} -m pip install matplotlib==3.7.3
# !{sys.executable} -m pip install imblearn
```

```
[69]: df = pd.read_csv('wine-data-set .csv')
df
```

```
[69]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.0	0.270	0.36	20.7	0.045	
1	6.3	0.300	0.34	1.6	0.049	
2	8.1	0.280	0.40	6.9	0.050	
3	7.2	0.230	0.32	8.5	0.058	
4	7.2	0.230	0.32	8.5	0.058	
...	...	...	...	...	...	
6458	6.8	0.620	0.08	1.9	0.068	
6459	6.2	0.600	0.08	2.0	0.090	
6460	6.3	0.510	0.13	2.3	0.076	
6461	5.9	0.645	0.12	2.0	0.075	
6462	6.0	0.310	0.47	3.6	0.067	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	45.0	170.0	1.00100	3.00	0.45	
1	14.0	132.0	0.99400	3.30	0.49	
2	30.0	97.0	0.99510	3.26	0.44	
3	47.0	186.0	0.99560	3.19	0.40	
4	47.0	186.0	0.99560	3.19	0.40	
...	...	...	...	...	...	
6458	28.0	38.0	0.99651	3.42	0.82	
6459	32.0	44.0	0.99490	3.45	0.58	
6460	29.0	40.0	0.99574	3.42	0.75	

6461	32.0	44.0	0.99547	3.57	0.71
6462	18.0	42.0	0.99549	3.39	0.66

	alcohol	quality
0	8.8	6
1	9.5	6
2	10.1	6
3	9.9	6
4	9.9	6
...	...	...
6458	9.5	6
6459	10.5	5
6460	11.0	6
6461	10.2	5
6462	11.0	6

[6463 rows x 12 columns]

[70]: df.describe()

```
[70]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	6463.000000	6463.000000	6463.000000	6463.000000
mean	7.217755	0.339589	0.318758	5.443958
std	1.297913	0.164639	0.145252	4.756852
min	3.800000	0.080000	0.000000	0.600000
25%	6.400000	0.230000	0.250000	1.800000
50%	7.000000	0.290000	0.310000	3.000000
75%	7.700000	0.400000	0.390000	8.100000
max	15.900000	1.580000	1.660000	65.800000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	6463.000000	6463.000000	6463.000000	6463.000000
mean	0.056056	30.516865	115.694492	0.994698
std	0.035076	17.758815	56.526736	0.003001
min	0.009000	1.000000	6.000000	0.987110
25%	0.038000	17.000000	77.000000	0.992330
50%	0.047000	29.000000	118.000000	0.994890
75%	0.065000	41.000000	156.000000	0.997000
max	0.611000	289.000000	440.000000	1.038980

	pH	sulphates	alcohol	quality
count	6463.000000	6463.000000	6463.000000	6463.000000
mean	3.218332	0.531150	10.492825	5.818505
std	0.160650	0.148913	1.193128	0.873286
min	2.720000	0.220000	8.000000	3.000000
25%	3.110000	0.430000	9.500000	5.000000
50%	3.210000	0.510000	10.300000	6.000000

75%	3.320000	0.600000	11.300000	6.000000
max	4.010000	2.000000	14.900000	9.000000

```
[71]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6463 entries, 0 to 6462
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          6463 non-null   float64
1   volatile acidity       6463 non-null   float64
2   citric acid            6463 non-null   float64
3   residual sugar         6463 non-null   float64
4   chlorides              6463 non-null   float64
5   free sulfur dioxide    6463 non-null   float64
6   total sulfur dioxide   6463 non-null   float64
7   density                6463 non-null   float64
8   pH                    6463 non-null   float64
9   sulphates              6463 non-null   float64
10  alcohol                6463 non-null   float64
11  quality                6463 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 606.0 KB
```

```
[72]: df.isna().sum()
```

```
[72]: fixed acidity          0
      volatile acidity     0
      citric acid          0
      residual sugar       0
      chlorides            0
      free sulfur dioxide   0
      total sulfur dioxide  0
      density              0
      pH                  0
      sulphates            0
      alcohol              0
      quality              0
      dtype: int64
```

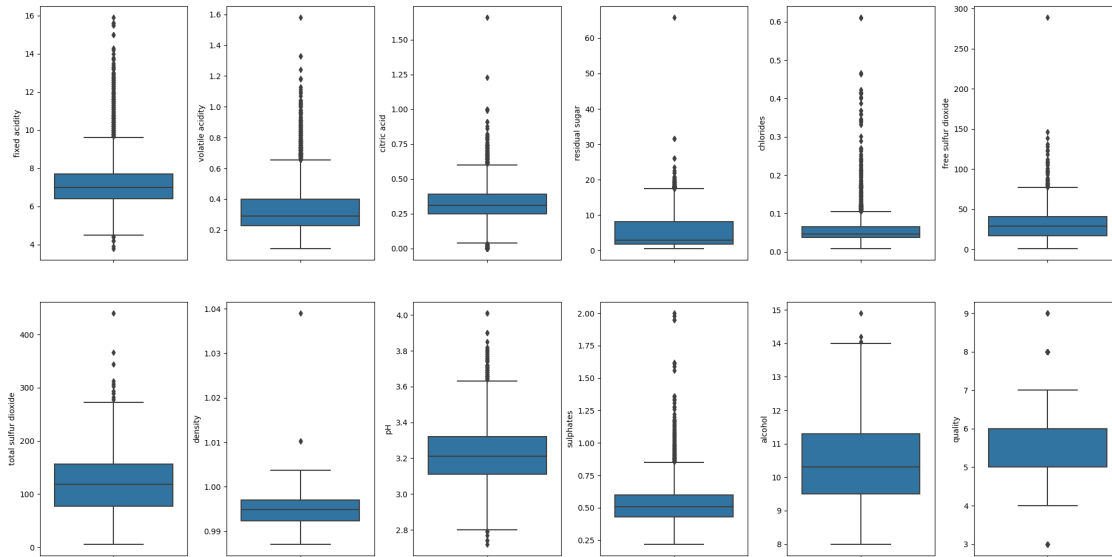
```
[73]: # create box plots
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in df.items():
    if col != 'type':
```

```

sns.boxplot(y=col, data=df, ax=ax[index])
index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)

```

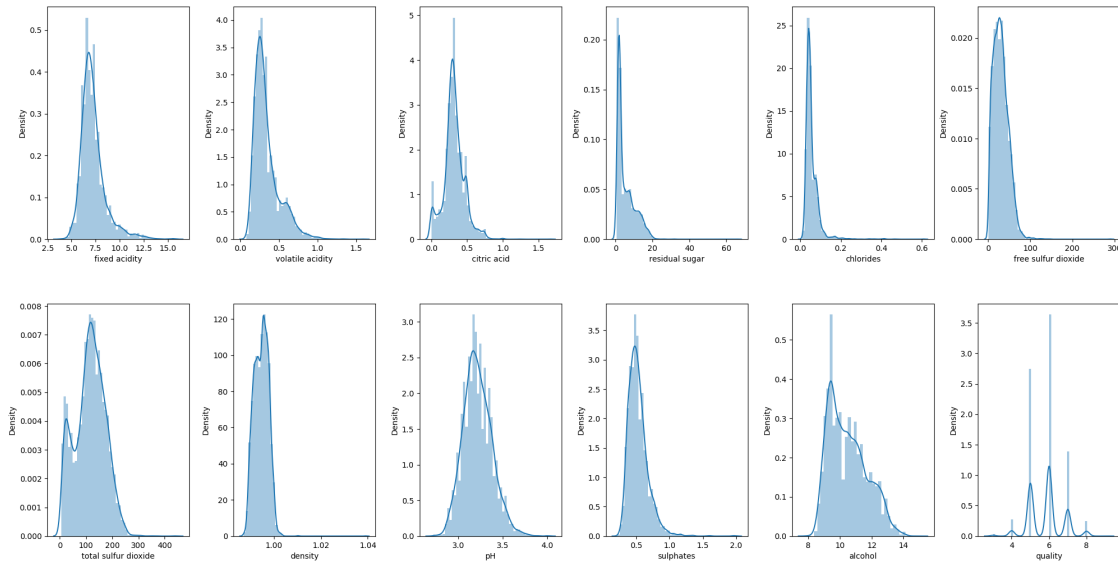


```

[74]: # create dist plot
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in df.items():
    if col != 'type':
        sns.distplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)

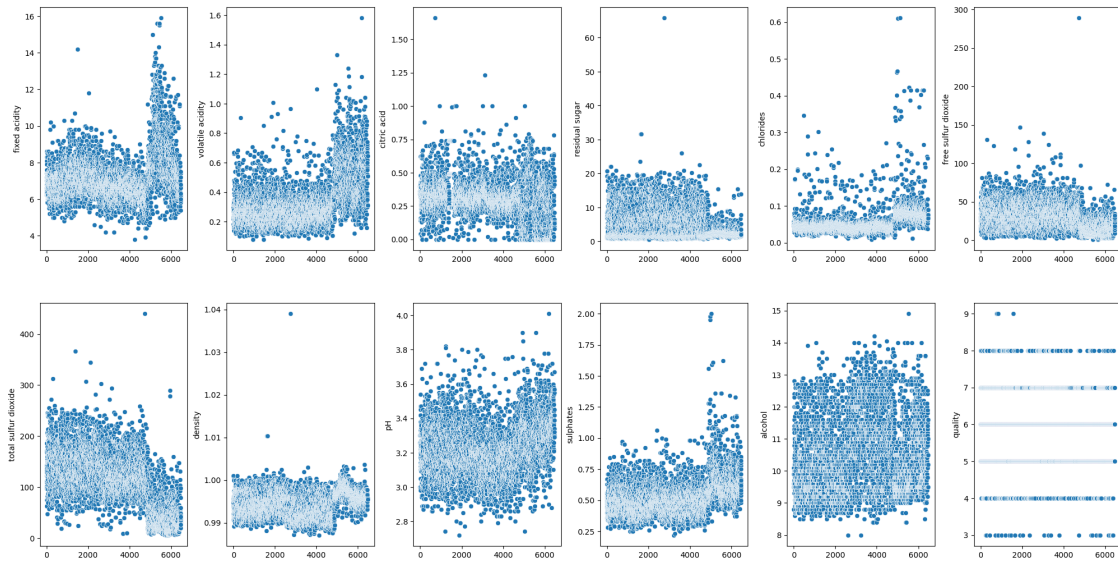
```



```
[75]: ## log transformation
      # df['free sulfur dioxide'] = np.log(1 + df['free sulfur dioxide'])
      # sns.distplot(df['free sulfur dioxide'])
```


```
[76]: # create dist plot
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in df.items():
    if col != 'type':
        sns.scatterplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[77]: sns.countplot(x= df['quality'])
```

```
[77]: <Axes: xlabel='quality', ylabel='count'>
```



```
[78]: corr = df.corr()
      corr
```

```
[78]:
```

	fixed acidity	volatile acidity	citric acid	\
fixed acidity	1.000000	0.221066	0.323744	
volatile acidity	0.221066	1.000000	-0.377512	
citric acid	0.323744	-0.377512	1.000000	
residual sugar	-0.113442	-0.196677	0.142324	
chlorides	0.299104	0.377995	0.039412	
free sulfur dioxide	-0.283485	-0.353402	0.132271	
total sulfur dioxide	-0.330543	-0.414729	0.194398	
density	0.459713	0.272101	0.097068	
pH	-0.251121	0.260134	-0.327860	
sulphates	0.301263	0.225656	0.059070	
alcohol	-0.096190	-0.039528	-0.010056	
quality	-0.076174	-0.266677	0.084926	

	residual sugar	chlorides	free sulfur dioxide	\
fixed acidity	-0.113442	0.299104	-0.283485	
volatile acidity	-0.196677	0.377995	-0.353402	
citric acid	0.142324	0.039412	0.132271	
residual sugar	1.000000	-0.128814	0.403449	
chlorides	-0.128814	1.000000	-0.195428	
free sulfur dioxide	0.403449	-0.195428	1.000000	
total sulfur dioxide	0.495684	-0.279602	0.721476	
density	0.551494	0.363108	0.025113	
pH	-0.266481	0.044653	-0.145164	
sulphates	-0.185616	0.396240	-0.188947	
alcohol	-0.359132	-0.257664	-0.179477	
quality	-0.034654	-0.200553	0.054924	


  

	total sulfur dioxide	density	pH	sulphates	\
fixed acidity	-0.330543	0.459713	-0.251121	0.301263	
volatile acidity	-0.414729	0.272101	0.260134	0.225656	
citric acid	0.194398	0.097068	-0.327860	0.059070	
residual sugar	0.495684	0.551494	-0.266481	-0.185616	
chlorides	-0.279602	0.363108	0.044653	0.396240	
free sulfur dioxide	0.721476	0.025113	-0.145164	-0.188947	
total sulfur dioxide	1.000000	0.031419	-0.237204	-0.275878	
density	0.031419	1.000000	0.012525	0.260019	
pH	-0.237204	0.012525	1.000000	0.190864	
sulphates	-0.275878	0.260019	0.190864	1.000000	
alcohol	-0.264385	-0.687432	0.120473	-0.004116	
quality	-0.041598	-0.304447	0.018403	0.039054	

	alcohol	quality
fixed acidity	-0.096190	-0.076174
volatile acidity	-0.039528	-0.266677
citric acid	-0.010056	0.084926
residual sugar	-0.359132	-0.034654
chlorides	-0.257664	-0.200553
free sulfur dioxide	-0.179477	0.054924
total sulfur dioxide	-0.264385	-0.041598
density	-0.687432	-0.304447
pH	0.120473	0.018403
sulphates	-0.004116	0.039054
alcohol	1.000000	0.444637
quality	0.444637	1.000000

```
[79]: plt.figure(figsize=(20,10))
sns.heatmap(corr, annot=True, cmap='coolwarm' )
```

[79]: <Axes: >



```
[80]: # Separate predictors (X) and target variable (y)
X = df.drop('quality', axis=1)
y = df['quality']
y.value_counts()
```



```
[80]: 6    2820
      5    2128
      7    1074
      4     214
      8     192
      3      30
      9       5
      Name: quality, dtype: int64
```

```
[81]: # classify function
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.metrics import accuracy_score
from imblearn.over_sampling import SMOTE

def classify(model, X, y):
    oversample = SMOTE(k_neighbors=4)
    # transform the dataset
    X, y = oversample.fit_resample(X, y)
    # print(y.value_counts())
    x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↪random_state=42)
    # train the model
    model.fit(x_train, y_train)

    # Predict on the test set
    y_pred = model.predict(x_test)
    bottom_line_accuracy = accuracy_score(y_test, y_pred)
    print(f"Bottom-Line Accuracy: {bottom_line_accuracy*100:.4f}%")
    # print("Accuracy:", model.score(x_test, y_test) * 100)

    # # cross-validation
    # score = cross_val_score(model, X, y, cv=5)
    # print("CV Score:", np.mean(score)*100)
```

```
[82]: from sklearn.linear_model import LogisticRegression

model = LogisticRegression(random_state=42)
print('Accuracy score without transforming any predictors:')
classify(model, X, y)
```

Accuracy score without transforming any predictors:  
Bottom-Line Accuracy: 31.3070%

```
[83]: # # log transformation
# df['free sulfur dioxide'] = np.log(1 + df['free sulfur dioxide'])
# sns.distplot(df['free sulfur dioxide'])
```

## 1 6. Applying Transformation on one predictor each time and getting accuracy

```
[84]: from sklearn.preprocessing import StandardScaler

scale_df = df.copy()
scaler = StandardScaler()
scaled_fixed_acidity = scaler.fit_transform(scale_df[['fixed acidity']])
scale_df['fixed acidity'] = scaled_fixed_acidity
X = df.drop('quality', axis=1)
y = df['quality']

print('Accuracy score with StandardScaler on predictor "fixed acidity":')
classify(model, X, y)
```

Accuracy score with StandardScaler on predictor "fixed acidity":  
Bottom-Line Accuracy: 31.9149%

```
[85]: from sklearn.preprocessing import MinMaxScaler

scale_df = df.copy()
scaler = MinMaxScaler()
scaled_alcohol = scaler.fit_transform(scale_df[['alcohol']])
scale_df['alcohol'] = scaled_alcohol
X = df.drop('quality', axis=1)
y = df['quality']

print('Accuracy score with MinMaxScaler on predictor "alcohol":')
classify(model, X, y)
```

Accuracy score with MinMaxScaler on predictor "alcohol":  
Bottom-Line Accuracy: 32.3202%

```
[86]: from sklearn.preprocessing import RobustScaler

scale_df = df.copy()
scaler = RobustScaler()
scaled_total_sulfur_dioxide = scaler.fit_transform(scale_df[['total sulfur_dioxide']])
scale_df['total sulfur dioxide'] = scaled_total_sulfur_dioxide
X = scale_df.drop('quality', axis=1)
y = scale_df['quality']

print('Accuracy score with RobustScaler on predictor "total sulfur dioxide":')
classify(model, X, y)
# X.describe()
```

Accuracy score with RobustScaler on predictor "total sulfur dioxide":  
Bottom-Line Accuracy: 36.9605%

```
[87]: from sklearn.preprocessing import MaxAbsScaler

scale_df = df.copy()
scaler = MaxAbsScaler()
scaled_free_sulfur_dioxide = scaler.fit_transform(scale_df[['free sulfur_
    dioxide']])
scale_df['free sulfur dioxide'] = scaled_free_sulfur_dioxide
X = scale_df.drop('quality', axis=1)
y = scale_df['quality']

# scaler = RobustScaler()
# X = scaler.fit_transform(X)

print('Accuracy score with MaxAbsScaler on predictor "free sulfur dioxide":')
classify(model, X, y)
# X.describe()
```

Accuracy score with MaxAbsScaler on predictor "free sulfur dioxide":  
Bottom-Line Accuracy: 31.5299%

```
[88]: from sklearn.preprocessing import QuantileTransformer

scale_df = df.copy()
scaler = QuantileTransformer()
scaled_residual_sugar = scaler.fit_transform(scale_df[['residual sugar']])
scale_df['residual sugar'] = scaled_residual_sugar

X = df.drop('quality', axis=1)
y = df['quality']

print('Accuracy score with QuantileTransformer on predictor "residual sugar":')
classify(model, X, y)
```

Accuracy score with QuantileTransformer on predictor "residual sugar":  
Bottom-Line Accuracy: 32.2999%

## 2 8. Selecting subsets and getting accuracy

```
[89]: # X = df.drop(columns=['quality', 'total sulfur dioxide', 'density'], axis=1)
X = df[['alcohol']]
y = df['quality']
# scaler = RobustScaler()
# X = scaler.fit_transform(X)
print('Selected subset ["alcohol"] based on correlation number closest to 1 for_
    "quality":')
```

```
classify(model, X, y)
```

Selected subset ["alcohol"] based on correlation number closest to 1 for "quality":

Bottom-Line Accuracy: 28.7741%

```
[90]: X = df[['alcohol', 'density', 'citric acid']]
y = df['quality']
# scaler = RobustScaler()
# X = scaler.fit_transform(X)
print("Selected subset [alcohol', 'density', 'citric acid'] on random:")
classify(model, X, y)
```

Selected subset [alcohol', 'density', 'citric acid'] on random:

Bottom-Line Accuracy: 30.1925%

```
[91]: X = df[['alcohol', 'citric acid', 'free sulfur dioxide']]
y = df['quality']
# scaler = RobustScaler()
# X = scaler.fit_transform(X)
print("Selected subset ['alcohol', 'citric acid', 'free sulfur dioxide'] based_
↳on correlation numbers closest to 1 for 'quality':")
classify(model, X, y)
```

Selected subset ['alcohol', 'citric acid', 'free sulfur dioxide'] based on correlation numbers closest to 1 for 'quality':

Bottom-Line Accuracy: 33.1307%

```
[92]: X = df[['density', 'volatile acidity', 'chlorides']]
y = df['quality']

print('Selected subset based on correlation numbers closest to negative 1 for_
↳"quality":')
classify(model, X, y)
```

Selected subset based on correlation numbers closest to negative 1 for "quality":

Bottom-Line Accuracy: 31.3070%

```
[93]: X = df[['alcohol', 'density', 'volatile acidity']]
y = df['quality']
# scaler = RobustScaler()
# X = scaler.fit_transform(X)
print('Selected subset based on correlation numbers closest to 1 or negative 1_
↳for "quality":')
classify(model, X, y)
```

Selected subset based on correlation numbers closest to 1 or negative 1 for "quality":

Bottom-Line Accuracy: 29.5846%

## Exploratory data analysis and initial model

To get The Bottom-Line Accuracy score of machine learning model named LogisticRegression I had to first look for null values in dataset. Then I looked for outliers using exploratory data analysis and plotting. I could have modified outlier data at this point in time. Since I was not asked to do it I moved on to the next step. Then I looked for correlation for all variables which later became useful for selecting subsets. After that I split the variable into X and y, where y is the target variable quality. For model training I created classify function which takes in model, X and y. since y had severely imbalanced classes I used oversampling with SMOTE on X,y to generate new features from minority classes. this made sure all the classes in y have oversampled to the upper value. I used train\_test\_split from sklearn to split X and y in Train and test datasets. I kept 25% of the data in test datasets and and rest in train dataset. I fit x\_train and y\_train into model. then I get the model to predict on x\_test to get y\_preds. then comparint y\_test and y\_preds using accuracy score function I get bottom-line accuracy score. Initially i use Logistic Recression model and passed it with X and y to classify function without transforming them.

Initial Bottom-Line Accuracy: 31.3070%

### Transforming on predictor each time

Accuracy score with StandardScaler on predictor “fixed acidity”: 31.9149%, Here Bottom-Line Accuracy increases from original prediction

Accuracy score with MinMaxScaler on predictor “alcohol”: 32.3202%, Here Bottom-Line Accuracy increases from original prediction

Accuracy score with RobustScaler on predictor “total sulfur dioxide”: 36.9605%, Here Bottom-Line Accuracy increases from original prediction

Accuracy score with MaxAbsScaler on predictor “free sulfur dioxide”: 31.5299%, Here Bottom-Line Accuracy increases from original prediction

Accuracy score with QuantileTransformer on predictor “residual sugar”: 32.2999%, Here Bottom-Line Accuracy increases from original prediction

So, I found that all of these transformations increased the accuracy score of the model. I also checked that if i used transformation of more than one predictor at each pass then accuracy improves further to close to 50% when RobustScaler is used on all predictors of X. Since this question restricts me to transform only one predictor at each pass I did not include it in the project.

### Selecting different subsets each time

Selected subset [“alcohol”] based on correlation number closest to 1 for “quality”: 28.7741%, Here Bottom-Line Accuracy decreases from original prediction

Selected subset [‘alcohol’, ‘density’, ‘citric acid’] on random: 30.1925%, Here Bottom-Line Accuracy decreases from original prediction

Selected subset [‘alcohol’, ‘citric acid’, ‘free sulfur dioxide’] based on correlation numbers closest to 1 for ‘quality’: 33.1307%, Here Bottom-Line Accuracy increases from original prediction

Selected subset [‘density’, ‘volatile acidity’, ‘chlorides’] based on correlation numbers closest to negative 1 for “quality”: 31.3070%, Here Bottom-Line Accuracy is same from original prediction

Selected subset ['alcohol', 'density', 'volatile acidity'] based on correlation numbers closest to 1 or negative 1 for "quality": 29.5846%, Here Bottom-Line Accuracy decreases from original prediction

So I found that only subset ['alcohol', 'citric acid', 'free sulfur dioxide'] increased the accuracy score of the model, Which is just as I expected since the 3 variables in this subset have the strongest positive correlation with quality variable. knowing the correlations was helpful to create this subset.

[ ]: